

# Machine Learning Performance Metrics

## Regression

Metric	Description	Pros	Cons
Mean Absolute Error (MAE)	Measures the average magnitude of errors between predictions and actual.	Simple to interpret; doesn't heavily penalize large errors.	Doesn't capture variability very well, less sensitive to outliers.
Mean Squared Error (MSE)	Measures average squared errors, penalizing large errors more than MAE.	Useful for algorithms that need error gradients (in optimization).	Sensitive to outliers, can disproportionately affect the metric.
Root Mean Squared Error (RMSE)	Square root of MSE, providing errors in same units as target variable.	Benefits of MSE but is more interpretable.	Sensitivity to outliers, as in MSE.
Mean Absolute Percentage Error (MAPE)	Expresses errors as a percentage of actual values.	Intuitive for business applications and relative comparisons.	Undefined when actual values are zero; biased towards underestimating errors.

<b>R-Squared (R2)</b>	Indicates the proportion of variance in the dependent variable explained by the model.	Easy to interpret, general indicator of model performance.	Can be misleading for small datasets or overfitting, doesn't penalize complexity.
<b>Adjust R-Squared</b>	Similar to R2, adjusts for the number of predictors, penalizing overfitting.	Accounts for model complexity; useful in feature selections.	Less intuitive to explain than R2.
<b>Akaike Information Criterion (AIC)</b>	Penalize model complexity while considering goodness of fit. Often used to compare separate models.	Useful for model comparison.	Not a standalone metric.

Classification

Metric	Description	Pros	Cons
Accuracy	Measures the proportion of correct predictions.	Easy to understand; broad indicator of performance.	Misleading for imbalanced datasets (high accuracy with majority class decision).
Precision	Proportion of true positives among predicted positives.	Useful when false positives are costly.	Doesn't account for false negatives.
Recall (Sensitivity)	Proportion of true positives identified out of all true positives.	Useful when false negatives are costly.	Doesn't account for false positives.
F1 Score	Harmonic mean of precision and recall.	Balances precision and recall, especially useful for imbalanced datasets.	Doesn't reflect true negatives; not as intuitive as individual metrics.
Specificity	Proportion of true negatives among actual negatives.	Complements recall; useful in medical diagnostics.	Not always emphasized in standard evaluations.
ROC/AUC	Plots true positive rate vs false positive rate at various thresholds.	Visualizes trade-offs between sensitivity and specificity.	

## Formulas

### Regression

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

$$MAPE = \frac{1}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right| \times 100$$

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

$$R_{adj}^2 = 1 - \frac{(1 - R^2)(n - 1)}{n - k - 1}$$

$$AIC = 2k - 2 \ln(L) \quad BIC = k \ln(n) - 2 \ln(L)$$

### Classification

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

$$Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP + FN}$$

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

$$Specificity = \frac{TN}{TN + FP}$$

$$TruePositiveRate = \frac{TP}{TP+FN} \text{ vs. } FalsePositiveRate = \frac{FP}{FP+TN}.$$